

## APPENDIX TO “ENABLING PERSONAL CONSENT IN DATABASES”, BY KONSTANTINIDIS ET. AL

PROOF. THEOREM 4.4 We need to prove that for all  $t \in q^\lambda(D)$  and support queries  $q_\pi = \pi_A(\hat{q}) \in \Pi_q$ , there is a constraint  $q_N$  and a tuple  $t_n \in q_N^\lambda(D)$  with  $\pi_A(\vec{t}) \sqcap t_n$  iff there is a minimal w.r.t. set inclusion  $S \subseteq \text{body}(q_\pi)$  with a query unifier of  $q_\pi$  over  $q_N$  for  $S$  having query unification  $q_m$ , and  $t_m \in q_m^\lambda(D)$  such that  $\pi_A(\vec{t}) \sqcap t_m$ .

$\Leftarrow$ ) First note that all query unifications, also the most general ones, of a query over a constraint are queries which are contained both in  $q_N$  and  $q_\pi$ . Moreover, for their query unifier  $\theta$ ,  $\theta$  is a containment mapping from the query and the constraint to the unification such that all annotated answer tuples of the unification will “propagate” the same annotations both to the query and the constraint, appropriately.

$\Rightarrow$ ) Consider the homomorphisms  $h_\pi$  from  $q_\pi$  to  $D$  and  $h_N$  from  $q_N$  to  $D$  that obtain  $\pi_A(\vec{t})$  and  $t_n$  respectively. Consider their composition and let database instance  $C$  (for “canonical”) the set of atoms in their image. Let  $\vec{x}, \vec{y}$  the head terms respectively of  $q_\pi, q_N$  (with  $|\vec{x}|=|\vec{y}|=|\vec{t}_n|=n$ ).

(Claim I): There is a minimal non-empty set of atoms  $S_m \subseteq C$  that contains all terms in  $\vec{t}_n$ , and a piece unifier  $\theta$  for a set  $S' \subseteq \text{body}(q_\pi)$  over a set  $S_N \subseteq \text{body}(q_N)$  such that  $\theta(S') \cup \theta(S_N) = S_m$  and for all atoms  $\alpha \in S_m$  there is a unifiable pair of  $\theta, (\beta, \delta)$ , and in at least one position that  $\alpha$  contains term  $\vec{t}_n[i]$ ,  $\beta$  and  $\delta$  contain head terms  $\vec{x}[i]$  and  $\vec{y}[i]$  respectively. We prove (Claim I) by contradiction. Suppose that for all  $S_m \subseteq C$  there exists an  $i \in [1, n]$  for which for some atom  $\alpha \in S_m$  with  $\vec{t}_n[i] \in \text{terms}(\alpha)$  there is no  $\beta \in \text{body}(q_\pi)$  with  $\vec{x}[i] \in \text{terms}(\beta)$ , and no  $\delta \in \text{body}(q_N)$  with  $\vec{y}[i] \in \text{terms}(\delta)$ , such that  $\theta(\{\beta, \delta\}) = \{\alpha\}$ . This means there are no mappings from  $q_\pi$  to  $S_m$  and from  $q_N$  to  $S_m$  that map two atoms, from  $q_\pi$  and  $q_N$  respectively, containing the  $i^{\text{th}}$  head term to the same atom  $\alpha$ ; if there were such mappings we could use their union (with no loss of generality we can assume  $q_\pi, q_N$  use disjoint variable names) as a unifier  $\theta$ . Since this is the case, on database  $C$  the  $i^{\text{th}}$  element of at least one tuple will be annotated with an identifier that will not be used for the  $i^{\text{th}}$  element of any resulting tuple of  $q_\pi$  nor the  $i^{\text{th}}$  element of any resulting tuple of  $q_N$  (since no atoms that contain the  $i^{\text{th}}$  element will map to our canonical database). However, this would prove that  $q_\pi$  and  $q_N$  do not “overlap” on  $C$  which is a contradiction, since, intuitively,  $C$  is just the part of  $D$  on which  $\pi_A(\vec{t})$  intersects with  $t_N$ . Having proved (Claim I), note that  $\theta$  is a most general query unifier of  $q_\pi$  over  $q_N$  for  $S'$ . We can union  $\theta$  with the restriction of  $h_\pi$  and  $h_N$  on the variables not in the domain of  $\theta$ , and obtain  $\theta_m$  which is still a most general unifier of  $q_\pi$  over  $q_N$  for  $S'$ . Note also that since  $q_m$  maps to  $C$  by construction it will return an annotated answer tuple  $t_m$  on  $C$  (and  $D$ ) that has the same base and the same schema as  $t_N$  with at least all annotation labels on each base term that  $t_N$  has on the same term. Thus it will be the case that  $\pi_A(\vec{t}) \sqcap t_m$ .  $\square$

Proof Sketch LEMMA 1 Notice that the outer SELECT statement executes the input query and the nested SELECT statement executes  $q_m$  without the conditions  $J_A$  and  $J_\theta$ . Moreover, notice that  $\vec{y}$  are terms from the input query ( $\pi_A(\hat{q})$ ) simply changes the head of the query by projecting terms and including constants), so by using  $J_A$  we recognize answer tuples  $\vec{t} \in q(D)$  and answer tuples  $\vec{t}_m \in q_m(D)$  for which  $\pi_A(\vec{t}) = \vec{t}_m$ , for any database  $D$ . The only additional step

that we need is to verify that the tuples that we remove, had we taken annotations into account, would be annotation-intersected with tuples  $t_m$  and those in  $q_N^\lambda(D)$ ; we do this via  $J_\theta$ , by enforcing that we remove tuples that came from the unified atoms, i.e. where all terms in the appropriate atoms  $\beta$  in the query are equated with terms in atoms  $\alpha$  in  $q_m$  ( $\alpha, \beta$  as in Thm 4.4).  $\square$

PROOF. THEOREM 4.5 (The problem is in NP) First notice that there is a polynomial number of most general query unifications for a query and a constraint. These are bounded by the initial pairs of query-constraint atoms that our algorithm tries to unify (the openECSets). Further unifications are tried only for combinations of non successfully unified pairs, and successful combinations are exploited further. Moreover, for performing the actual unifications our algorithm goes over polynomially many equivalence classes trying to merge those that have common elements (each class has polynomially many elements). Hence, we need to check that  $\vec{t}$  satisfies the queries of lemma 1 and overall we can check this in polynomial time. Thus the problem is in NP. (The problem is NP-hard) Deciding whether  $\vec{t}$  is consent-abiding implies deciding whether  $\vec{t}$  is an answer of the query to begin with (with no negative constraints at all) which is NP-hard[4].  $\square$

Proof Sketch When can easily prove the contrapositive: if the constraint is not secure with respect to the query, there is a critical tuple shared between the queries, and this means there are homomorphisms from both queries to the same actual tuple. Therefore, there will be an (empty) answer tuple in the answer of both queries annotated with (at least) the same identifier coming from the critical tuple. Thus, the query will violate the constraint.  $\square$